

CT5102: Programming for Data Analytics

4. Matrices and Data Frames

Prof. Jim Duggan,
School of Computer Science
University of Galway.

https://github.com/JimDuggan/explore_or

5	Matrices and Data Frames	83
5.1	Introduction	83
5.2	Matrices	84
5.3	Data frames	91
5.4	R functions for processing data frames: subset() and transform()	94
5.5	Tibbles	95
5.6	Functionals on matrices and data frames	97
5.6.1	Using apply() on a matrix and data frame	97
5.6.2	Using lapply() on data frames	99
5.6.3	Adding extra arguments to apply() and lapply() . . .	100
5.7	Mini-case 1: Modelling social networks using matrices	101
5.8	Mini-case 2: Creating a pipeline for processing data frames . .	103
5.9	Summary of R functions from Chapter 5	104
5.10	Exercises	105

On an intuitive level, a *data frame* is like a matrix, with a two-dimensional rows-and-columns structure. However, it differs from a matrix in that each column may have a different type.

— Norman Matloff ([Matloff, 2011](#))

Overview

- To date, we have used atomic vectors and lists to store information.
- While these are foundational data structures in R, they do not provide support for processing rectangular data, which is a common format in data science.
- More generally, we can have two types of rectangular (two-dimensional) data of :
 1. The same type, typically numeric, that is stored in a **matrix**, and
 2. Different types that is stored in a **data frame**.

(4.1) Matrices

- a matrix is a two-dimensional structure, with rows and columns, that contains the same type.
- it is an atomic vector in two dimensions, and is created using the `matrix()` function, with the following arguments:
 - `data`, which are the initial values, contained in an atomic vector, supplied to the matrix,
 - `nrow`, the desired number of rows,
 - `ncol`, the desired number of columns,
 - `byrow`, a logical value (default is FALSE), that specifies what way to fill the matrix with data, either filled by row or by column,
 - `dimnames`, a list of length 2 giving row and column names, respectively.

```

set.seed(100)
data <- sample(1:9)
data
#> [1] 7 6 3 1 2 5 9 4 8
m1 <- matrix(data,
             nrow = 3,
             ncol=3,
             dimnames = list(c("R1","R2","R3"),
                             c("C1","C2","C3")))

# Display the matrix
m1
#>      C1 C2 C3
#> R1   7  1  9
#> R2   6  2  4
#> R3   3  5  8

```

```

# Show the number of rows and columns
nrow(m1)
#> [1] 3
ncol(m1)
#> [1] 3

# Show the matrix dimensions
dim(m1)
#> [1] 3 3

```

Useful points

- The matrix is populated by column order as the default.
- If `by_row` was set to TRUE, then the matrix would be populated by row order.
- The row and column names are set using the `dimnames` argument. This is not required, and row names and column names can always be set on a matrix using the functions `rownames()` and `colnames()`.
- The functions `nrow()` and `ncol()` can be used to return the matrix dimensions.
- The function `dim()` provides information on the matrix dimensions, and can also be used to resize a matrix, for example, converting a 3×3 to a 1×9 .

Growing a matrix

- An important property of a matrix is that it can be extended
 - Add rows using `rbind()`
 - Add columns with `cbind()`

```
m1_r <- rbind(m1,c(1,2,3))
rownames(m1_r)[4] <- "R4"
m1_r
#>      C1 C2 C3
#> R1    7  1  9
#> R2    6  2  4
#> R3    3  5  8
#> R4    1  2  3
```

```
m1_c <- cbind(m1_r,c(10,20,30,40))
colnames(m1_c)[4] <- "C4"
m1_c
#>      C1 C2 C3 C4
#> R1    7  1  9 10
#> R2    6  2  4 20
#> R3    3  5  8 30
#> R4    1  2  3 40
```


Subsetting a matrix, like vectors with 2 dimensions

```
m1
#>      C1 C2 C3
#> R1   7  1  9
#> R2   6  2  4
#> R3   3  5  8

# Extract the value in row 2, column 2
m1 [2,2]
#> [1] 2
```

```
m1[1:2,1:2]
#>      C1 C2
#> R1   7  1
#> R2   6  2
```

```
# Extract first row and all columns
m1[1,]
#> C1 C2 C3
#>  7  1  9
```

```
# Extract first column, returned as a vector
```

```
m1[,1]
```

```
#> R1 R2 R3
```

```
#> 7 6 3
```

```
# Extract all rows and the first column, returned as matrix
```

```
m1[,1,drop=FALSE]
```

```
#> C1
```

```
#> R1 7
```

```
#> R2 6
```

```
#> R3 3
```

```
# Extract first row and first two columns
```

```
m1["R1",c("C1","C2")]
```

```
#> C1 C2
```

```
#> 7 1
```

```
# Extract all rows and first two columns
```

```
m1[,c("C1","C2")]
```

```
#> C1 C2
```

```
#> R1 7 1
```

```
#> R2 6 2
```

```
#> R3 3 5
```

Using logical vectors and `is.matrix()`

```
m1
#>      C1 C2 C3
#> R1    7  1  9
#> R2    6  2  4
#> R3    3  5  8
m1[c(T,F),]
#>      C1 C2 C3
#> R1    7  1  9
#> R3    3  5  8
```

```
A <- matrix(1:4,nrow=2)
B <- matrix(1:4,nrow=2,byrow = T)
C <- list(c1=1:2, c2=3:4)
is.matrix(A)
#> [1] TRUE
is.matrix(B)
#> [1] TRUE
is.matrix(C)
#> [1] FALSE
```

Arithmetic Operators – Element-wide basis

A

```
#>      A_C1 A_C2
#> A_R1     1   3
#> A_R2     2   4
```

B

```
#>      B_C1 B_C2
#> B_R1     1   2
#> B_R2     3   4
```

Multiplication of A and B

A*B

```
#>      A_C1 A_C2
#> A_R1     1   6
#> A_R2     6  16
```

Addition of A and B

A+B

```
#>      A_C1 A_C2
#> A_R1     2   5
#> A_R2     5   8
```

Other useful matrix functions

```
# Use matrix algebra to multiply two matrices
```

```
A%*%B
```

```
#>      B_C1 B_C2
```

```
#> A_R1    10    14
```

```
#> A_R2    14    20
```

```
t(A)
```

```
#>      A_R1 A_R2
```

```
#> A_C1     1     2
```

```
#> A_C2     3     4
```

```
rownames(A)
```

```
#> [1] "A_R1" "A_R2"
```

```
colnames(A)
```

```
#> [1] "A_C1" "A_C2"
```

```
dim(A)
```

```
#> [1] 2 2
```

```
dimnames(A)
```

```
#> [[1]]
```

```
#> [1] "A_R1" "A_R2"
```

```
#>
```

```
#> [[2]]
```

```
#> [1] "A_C1" "A_C2"
```

```
rowSums(A)
```

```
#> A_R1 A_R2
```

```
#>      4     6
```

```
rowMeans(A)
```

```
#> A_R1 A_R2
```

```
#>      2     3
```

```
colSums(A)
```

```
#> A_C1 A_C2
```

```
#>      3     7
```

```
colMeans(A)
```

```
#> A_C1 A_C2
```

```
#> 1.5 3.5
```

Matrices - Summary

- R provides good support for problems that require matrix manipulation, *but all values need to be the same type*
- Matrices need to be defined using the `matrix()` function.
- Many of the subsetting commands used for atomic vectors can also be used for matrices, and that includes referencing elements by the row/column name.
- Matrices can be extended easily, using functions such as `cbind()` and `rbind()`.

(4.2) Data Frames

- A data frame is similar to a matrix, with a two-dimensional row and column structure, while on a technical level, **a data frame is a list**, with the elements of that list containing equal length vectors (Matloff, 2011).
- It's defined using the **data.frame()** function
- The elements (columns) of a data frame can be of different types
- The data frame, with its row and column structure, will be familiar to anyone who has used a spreadsheet, where each column is a variable, and every row is an observation.
- The data frame, and its successor, the **tibble**, will be used extensively during this course

```
d <- data.frame(Number=1:5,
                 Letter=LETTERS[1:5],
                 Flag=c(T,F,T,F,T),
                 stringsAsFactors = F)
```

```
d
#>   Number Letter Flag
#> 1      1      A  TRUE
#> 2      2      B FALSE
#> 3      3      C  TRUE
#> 4      4      D FALSE
#> 5      5      E  TRUE
```

```
summary(d)
```

```
#>      Number      Letter      Flag
#>  Min.    :1  Length:5      Mode :logical
#> 1st Qu.:2   Class :character FALSE:2
#> Median :3   Mode  :character  TRUE :3
#> Mean    :3
#> 3rd Qu.:4
#> Max.    :5
```


Activities on a data frame

- An important activity that is required with a data frame is to be able to: (1) subset rows, (2) subset columns, and (3) add new columns.
- Because a data frame is a **list** and also shares properties of a matrix, we can combine subsetting mechanisms from both of these data structures to subset a data frame
- We can access a data frame column using the \$ operator.

Subsetting examples

d

```
#>   Number Letter  Flag
#> 1      1      A  TRUE
#> 2      2      B FALSE
#> 3      3      C  TRUE
#> 4      4      D FALSE
#> 5      5      E  TRUE
```

```
d[d$Flag == T,]
```

```
#>   Number Letter  Flag
#> 1      1      A  TRUE
#> 3      3      C  TRUE
#> 5      5      E  TRUE
```

```
d[1:2,]
```

```
#>   Number Letter  Flag
#> 1      1      A  TRUE
#> 2      2      B FALSE
```

```
d[1:2,c("Letter","Flag")]
```

```
#>   Letter  Flag
#> 1      A  TRUE
#> 2      B FALSE
```

Adding a new column

```
d1 <- d
d1$letter <- letters[1:5]
```

d1

```
#>      Number Letter  Flag letter
#> 1         1      A  TRUE      a
#> 2         2      B FALSE      b
#> 3         3      C  TRUE      c
#> 4         4      D FALSE      d
#> 5         5      E  TRUE      e
```

```
d2 <- cbind(d, letter2=letters[6:10])
```

d2

```
#>      Number Letter  Flag letter2
#> 1         1      A  TRUE      f
#> 2         2      B FALSE      g
#> 3         3      C  TRUE      h
#> 4         4      D FALSE      i
#> 5         5      E  TRUE      j
```

The `subset()` function

- The function `subset(x, subset, select)` returns subsets of vectors, matrices, or data frames that meet specified conditions.
- The main arguments to provide when subsetting data frames are:
 - `x`, the object to be subsetted,
 - `subset`, a logical expression indicating which rows should be kept,
 - `select`, which indicates the columns to be selected from the data frame. If this is not present, all columns are returned.

Examples (and alternatives)

```
subset(mtcars,mpg>32,select=c("mpg","disp"))  
#>           mpg disp  
#> Fiat 128      32.4 78.7  
#> Toyota Corolla 33.9 71.1
```

```
mtcars[mtcars[, "mpg"]>32,c("mpg","disp")]  
#>           mpg disp  
#> Fiat 128      32.4 78.7  
#> Toyota Corolla 33.9 71.1  
mtcars[mtcars$mpg>32,c("mpg","disp")]  
#>           mpg disp  
#> Fiat 128      32.4 78.7  
#> Toyota Corolla 33.9 71.1
```

The `transform()` function

- A second function that can be used to manipulate data frames is `transform(data, ...)`, which takes in the following arguments:
 - `data`, which is the data frame,
 - `...` which are additional arguments that capture the details of how the new column is created.

```
df1 <- subset(mtcars,mpg>32,select=c("mpg","disp"))
df1 <- transform(df1,kpg=mpg*1.6)
df1
#>           mpg disp  kpg
#> Fiat 128      32.4 78.7 51.84
#> Toyota Corolla 33.9 71.1 54.24
```

```
df1 <- subset(mtcars,mpg>32,select=c("mpg","disp"))
df1$kpg <- df1$mpg*1.6
df1
#>           mpg disp  kpg
#> Fiat 128      32.4 78.7 51.84
#> Toyota Corolla 33.9 71.1 54.24
```

(4.4) Tibbles

- Tibbles are a type of data frame; however they alter some data frame behaviors:
 - Printing, where tibbles only show the first ten rows, and limit the visible columns to those that fit on the screen. The type is also displayed for each variable.
 - Subsetting, where a tibble is always returned, and also partial matching is not supported.

Previous example – using a tibble

```
library(tibble)
d1 <- tibble(Number=1:5,
              Letter=LETTERS[1:5],
              Flag=c(T,F,T,F,T))
```

```
d1
#> # A tibble: 5 x 3
#>   Number Letter Flag
#>   <int> <chr>   <lgl>
#> 1     1 A      TRUE
#> 2     2 B     FALSE
#> 3     3 C      TRUE
#> 4     4 D     FALSE
#> 5     5 E      TRUE
```


Differences with data.frame

```
# Show the data frame
```

```
str(d)
```

```
#> 'data.frame':    5 obs. of  3 variables:
```

```
#> $ Number: int  1 2 3 4 5
```

```
#> $ Letter: chr  "A" "B" "C" "D" ...
```

```
#> $ Flag : logi  TRUE FALSE TRUE FALSE TRUE
```

```
# Show the tibble
```

```
str(d1)
```

```
#> tibble [5 x 3] (S3: tbl_df/tbl/data.frame)
```

```
#> $ Number: int [1:5] 1 2 3 4 5
```

```
#> $ Letter: chr [1:5] "A" "B" "C" "D" ...
```

```
#> $ Flag : logi [1:5] TRUE FALSE TRUE FALSE TRUE
```

Subsetting differences

```
# Subset the data frame
d[1:2,"Letter"]
#> [1] "A" "B"

# Subset the tibble
d1[1:2,"Letter"]
#> # A tibble: 2 x 1
#>   Letter
#>   <chr>
#> 1 A
#> 2 B
```

Moving between two types

```
str(as_tibble(d))
```

```
#> tibble [5 x 3] (S3: tbl_df/tbl/data.frame)
```

```
#> $ Number: int [1:5] 1 2 3 4 5
```

```
#> $ Letter: chr [1:5] "A" "B" "C" "D" ...
```

```
#> $ Flag : logi [1:5] TRUE FALSE TRUE FALSE TRUE
```

```
str(as.data.frame(d1))
```

```
#> 'data.frame': 5 obs. of 3 variables:
```

```
#> $ Number: int 1 2 3 4 5
```

```
#> $ Letter: chr "A" "B" "C" "D" ...
```

```
#> $ Flag : logi TRUE FALSE TRUE FALSE TRUE
```

(4.5) Functionals on matrices and data frames

- The `apply(x,margin,f)` function (similar to `lapply()`) is a functional used to iterate over matrices and data frames, and it accepts the following arguments:
 - `x`, which can be a matrix or a data frame.
 - `margin`, a number that indicates whether the iteration is by row (`margin=1`), or by column (`margin=2`).
 - `f`, which is the function to be applied during each iteration.

Using `apply()` – analysing grades

```
set.seed(100)
grades <- sample(30:90,15,replace = T)
results <- matrix(grades,nrow=5)
rownames(results) <- paste0("St-",1:5)
colnames(results) <- paste0("Sub-",1:3)
results
```

```
#>      Sub-1 Sub-2 Sub-3
#> St-1     39    54    51
#> St-2     84    87    35
#> St-3     67    43    33
#> St-4     77    73    84
#> St-5     80    52    35
```

Maximum Grade for each subject

```
results
#>      Sub-1 Sub-2 Sub-3
#> St-1     39    54    51
#> St-2     84    87    35
#> St-3     67    43    33
#> St-4     77    73    84
#> St-5     80    52    35
```

```
max_gr_subject <- apply(results,           # the matrix
                        2,                 # 2 for columns
                        function(x)max(x)) # the function to apply

max_gr_subject
#> Sub-1 Sub-2 Sub-3
#>    84    87    84
```

Maximum Grade for each student

```
results
#>      Sub-1 Sub-2 Sub-3
#> St-1     39   54   51
#> St-2     84   87   35
#> St-3     67   43   33
#> St-4     77   73   84
#> St-5     80   52   35
```

```
max_gr_student <- apply(results,           # the matrix
                        1,                 # 1 for rows
                        function(x)max(x)) # the function to apply
```

```
max_gr_student
```

```
#> St-1 St-2 St-3 St-4 St-5
#>   54   87   67   84   80
```

Using `apply()` on data frames

```
set.seed(100)
my_mtcars <- mtcars[sample(1:6),c("mpg","cyl","disp")]
rows <- sample(1:nrow(my_mtcars),5)
rows
#> [1] 6 4 3 2 5

my_mtcars[rows[1],1] <- NA
my_mtcars[rows[2],2] <- NA
my_mtcars[rows[3],3] <- NA
```

```
my_mtcars[rows[4],1] <- NA
my_mtcars[rows[5],2] <- NA
my_mtcars
#>                mpg  cyl  disp
#> Mazda RX4 Wag    21.0    6   160
#> Datsun 710        NA     4   108
#> Mazda RX4        21.0    6    NA
#> Valiant           18.1   NA   225
#> Hornet Sportabout 18.7   NA   360
#> Hornet 4 Drive    NA     6   258
```


Count number of missing values by row

```
#>           mpg cyl disp
#> Mazda RX4 Wag      21.0   6  160
#> Datsun 710         NA    4  108
#> Mazda RX4         21.0   6   NA
#> Valiant           18.1  NA  225
#> Hornet Sportabout  18.7  NA  360
#> Hornet 4 Drive     NA    6  258
```

```
n_rm <- apply(my_mtcars,1,function(x)sum(is.na(x)))
n_rm
#>      Mazda RX4 Wag      Datsun 710      Mazda RX4
#>              0              1              1
#>      Valiant Hornet Sportabout      Hornet 4 Drive
#>              1              1              1
sum(n_rm)
#> [1] 5
```

Count number of missing values by column

```
#>           mpg cyl disp
#> Mazda RX4 Wag      21.0   6  160
#> Datsun 710          NA    4  108
#> Mazda RX4          21.0   6   NA
#> Valiant            18.1  NA  225
#> Hornet Sportabout  18.7  NA  360
#> Hornet 4 Drive     NA    6  258
```

```
n_cm <- apply(my_mtcars,2,function(x)sum(is.na(x)))
n_cm
#>   mpg   cyl  disp
#>    2     2     1
sum(n_cm)
#> [1] 5
```

Using `lapply()` on data frames

- Given that a data frame is also a list, and that `lapply()` processes lists, it also means that the `lapply()` functional can be used to process a data frame.
- When processing data frames with `lapply()`, the most important thing to remember is that the data frame will be processed *column-by-column* (e.g. variable)

```
str(mtcars)
#> 'data.frame':   32 obs. of  11 variables:
#> $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.
#> $ cyl : num  6 6 4 6 8 6 8 4 4 6 ...
#> $ disp: num  160 160 108 258 360 ...
#> $ hp  : num  110 110 93 110 175 105 245 62
#> $ drat: num  3.9 3.9 3.85 3.08 3.15 2.76 3
#> $ wt  : num  2.62 2.88 2.32 3.21 3.44 ...
#> $ qsec: num  16.5 17 18.6 19.4 17 ...
#> $ vs  : num  0 0 1 1 0 1 0 1 1 1 ...
#> $ am  : num  1 1 1 0 0 0 0 0 0 0 ...
#> $ gear: num  4 4 4 3 3 3 3 4 4 4 ...
#> $ carb: num  4 4 1 1 2 1 4 2 2 4 ...
```

```
s1 <- mtcars |>
  subset(select=c("mpg","cyl","disp")) |>
  lapply(function(x)mean(x))

s1
#> $mpg
#> [1] 20.09
#>
#> $cyl
#> [1] 6.188
#>
#> $disp
#> [1] 230.7
```

(4.6) Mini-case: Creating a pipeline for processing data frames

- For the data frame `mtcars`, the following processing actions will be taken:
 - Two columns from `mtcars` will be selected, `mpg` and `disp`.
 - A new column `kpg` will be added, which converts `mpg` to kilometers per gallon, using the multiplier `1.6`.
 - A new column `dm_ratio` will be added, which is the ratio of `disp` and `mpg`.
 - The first six observations will then be shown.

Solution pipeline

```
mtcars_1 <- mtcars |> # the original data frame
  subset(select=c("mpg","disp")) |> # select 2 columns
  transform(kpg=mpg*1.6, # Add first column
            dm_ratio=disp/mpg) |> # Add second column
  head() # Subset 1st 6 records
```

mtcars_1

```
#>      mpg disp  kpg dm_ratio
#> Mazda RX4      21.0  160 33.60      7.619
#> Mazda RX4 Wag  21.0  160 33.60      7.619
#> Datsun 710     22.8  108 36.48      4.737
#> Hornet 4 Drive  21.4  258 34.24     12.056
#> Hornet Sportabout 18.7  360 29.92     19.251
#> Valiant        18.1  225 28.96     12.431
```

(4.7) Summary Functions

Function	Description
<code>as.data.frame()</code>	Converts a tibble to a data frame
<code>apply()</code>	Iterates over rectangular data, by row or by column.
<code>cbind()</code>	Adds a new vector as a matrix column.
<code>colnames()</code>	Set (or view) the column names of a matrix.
<code>colMeans()</code>	Calculates the mean of each column in a matrix.
<code>colSums()</code>	Calculates the sum of each column in a matrix.
<code>data.frame()</code>	Constructs a data frame.
<code>diag()</code>	Sets a matrix diagonal, or generates an identity matrix
<code>dim()</code>	Returns (or sets) the matrix dimensions.
<code>dimnames()</code>	Returns the row and column names of a matrix.

<code>eigen()</code>	Calculates matrix eigenvalues and eigenvectors.
<code>factor()</code>	Encode a vector as a factor.
<code>is.matrix()</code>	Checks to see if the object is a matrix.
<code>matrix()</code>	Creates a matrix from the given set of arguments.
<code>rbind()</code>	Adds a vector as a row to a matrix.
<code>rownames()</code>	Sets (or views) the row names of a matrix.
<code>t()</code>	Returns the matrix transpose.
<code>rowMeans()</code>	Calculates the mean of each matrix row.
<code>rowSums()</code>	Calculates the sum of each matrix row.
<code>subset()</code>	Subsets data frames which meet specified conditions.
<code>tibble()</code>	Constructs a tibble (tibble package).
<code>as_tibble()</code>	Converts a data frame to a tibble (tibble package).
<code>transform()</code>	Add columns to a data frame.

(4.8) Exercises

1. Use the following initial code to generate the matrix `res`.

```
set.seed(100)
N=10
CX101 <- rnorm(N,45,8)
CX102 <- rnorm(N,65,8)
CX103 <- rnorm(N,85,25)
CX104 <- rnorm(N,60,15)
CX105 <- rnorm(N,55,15)
```

```
res
#>      CX101 CX102 CX103 CX104 CX105
#> Student-1 40.98 65.72  74.05 58.63 53.48
#> Student-2 46.05 65.77 104.10 86.36 76.05
#> Student-3 44.37 63.39  91.55 57.93 28.35
#> Student-4 52.09 70.92 104.34 58.33 64.34
#> Student-5 45.94 65.99  64.64 49.65 47.17
#> Student-6 47.55 64.77  74.04 56.67 74.83
#> Student-7 40.35 61.89  66.99 62.74 49.55
#> Student-8 50.72 69.09  90.77 66.26 74.79
#> Student-9 38.40 57.69  56.06 75.98 55.66
#> Student-10 42.12 83.48  91.18 74.55 26.82
```

2. The matrix `res` (from the previous question) has values that are out of the valid range for grades (i.e., greater than 100). To address this, all out-of-range values should be replaced by `NA`. Use `apply()` to generate the following modified matrix.

```
res_clean
```

```
#>           CX101 CX102 CX103 CX104 CX105
#> Student-1  40.98 65.72 74.05 58.63 53.48
#> Student-2  46.05 65.77    NA 86.36 76.05
#> Student-3  44.37 63.39 91.55 57.93 28.35
#> Student-4  52.09 70.92    NA 58.33 64.34
#> Student-5  45.94 65.99 64.64 49.65 47.17
#> Student-6  47.55 64.77 74.04 56.67 74.83
#> Student-7  40.35 61.89 66.99 62.74 49.55
#> Student-8  50.72 69.09 90.77 66.26 74.79
#> Student-9  38.40 57.69 56.06 75.98 55.66
#> Student-10 42.12 83.48 91.18 74.55 26.82
```

3. The matrix `res_clean` (from the previous question) has NA values, and as a work around, it has been decided to replace these values with the average subject mark. Write the code (using `apply()`) to generate the matrix `res_update`.

```
res_update
#>           CX101 CX102 CX103 CX104 CX105
#> Student-1  40.98 65.72 74.05 58.63 53.48
#> Student-2  46.05 65.77 76.16 86.36 76.05
#> Student-3  44.37 63.39 91.55 57.93 28.35
#> Student-4  52.09 70.92 76.16 58.33 64.34
#> Student-5  45.94 65.99 64.64 49.65 47.17
#> Student-6  47.55 64.77 74.04 56.67 74.83
#> Student-7  40.35 61.89 66.99 62.74 49.55
#> Student-8  50.72 69.09 90.77 66.26 74.79
#> Student-9  38.40 57.69 56.06 75.98 55.66
#> Student-10 42.12 83.48 91.18 74.55 26.82
```

4. Use the `subset()` function to generate the following tibbles from the tibble `ggplot2::mpg`. Use the R pipe operator (`|>`) where necessary.

```
# The car with the maximum displacement, with a subset of features
max_displ
```

```
#> # A tibble: 1 x 6
```

```
#>   manufacturer model      year displ  cty class
#>   <chr>          <chr>    <int> <dbl> <int> <chr>
#> 1 chevrolet    corvette  2008     7    15 2seater
```

```
# All 2seater cars, with selected columns
```

```
two_seater
```

```
#> # A tibble: 5 x 6
```

```
#>   class      manufacturer model      displ  year  hwy
#>   <chr>    <chr>          <chr>    <dbl> <int> <int>
#> 1 2seater chevrolet    corvette  5.7  1999  26
#> 2 2seater chevrolet    corvette  5.7  1999  23
#> 3 2seater chevrolet    corvette  6.2  2008  26
#> 4 2seater chevrolet    corvette  6.2  2008  25
#> 5 2seater chevrolet    corvette  7    2008  24
```