


Data Science for Operational Researchers Using R Online

6. Relational data with dplyr

Prof. Jim Duggan,
School of Computer Science
University of Galway.

https://github.com/JimDuggan/explore_or


Day 2 Overview




It's rare that a data analysis involves only a single table of data.
Typically you have many tables of data, and you must combine
them to answer the questions you are interested in.

— Hadley Wickham and Garrett Grolemund ([Wickham and Grolemund, 2016](#))

Overview



1. Overview – relational data




Leveraging data relationships

- When exploring data, there may be more than one tibble of interest, and these tibbles may share common information
- For example, we may want to explore associations between:
 - flight delays and weather events, and these information sources may be contained in two different tibbles.
 - weather data and wind energy generated, in order to build energy prediction models.

Relational concepts

- A relational model organizes data into one or more tables of columns and rows, where a unique key identifies each row.
- A key can be a single column value, known as a **primary key**, for example, a unique identifier for a person.
- For example, for the tibble `aimsir17::stations`, the primary key is the variable station, as this value is unique for each row, and therefore can be used to identify a single observation.
- A primary key from one table can also be a column in another table, and if this is the case, it is termed a **foreign key** in that table.

Order contains one product, product can be part of zero or more orders



Some observations

1. To simplify the example, we limit the number of product types for each order to one. Therefore, we have what is called a one-to-many relationship, where an order is for one product, but a product can be part of many orders.
2. There is an in-built inconsistency in the data, specifically, one of the orders (“OR-4”) contains a product (“PR-4”) that is not present in the product table. Normally, in a real-world database this would not be allowed, and through a process known as referential integrity, the database management system would enforce a constraint that a foreign key in one table would have to already be defined as a primary key in another table.
3. We use (2) to show a range of joins available in dplyr.

Creating the data - products

```
products <- tibble(ProductID=c("PR-1","PR-2","PR-3","PR-9"),  
                      Type=c("Computer","Tablet","Phone","Headphones"))  
  
products  
#> # A tibble: 4 x 2  
#>   ProductID  Type  
#>   <chr>       <chr>  
#> 1 PR-1        Computer  
#> 2 PR-2        Tablet  
#> 3 PR-3        Phone  
#> 4 PR-9        Headphones
```

Creating the data - orders

```
orders <- tibble(OrderID=c("OR-1","OR-2","OR-3","OR-4","OR-5"),  
                  Quantity=c(1,2,1,2,3),  
                  ProductID=c("PR-1","PR-2","PR-3","PR-4","PR-1"))  
  
orders  
#> # A tibble: 5 x 3  
#>   OrderID  Quantity ProductID  
#>   <chr>     <dbl>   <chr>  
#> 1 OR-1       1 PR-1  
#> 2 OR-2       2 PR-2  
#> 3 OR-3       1 PR-3  
#> 4 OR-4       2 PR-4  
#> 5 OR-5       3 PR-1
```

2. Mutating joins

- The first category of joins we explore is known as a mutating join, as these result in a new tibble with **additional columns**.
- Here are dplyr functions that perform mutating joins:
 - `inner_join()`,
 - `left_join()`,
 - `full_join()`.

`inner_join(x,y)`

- The `inner_join()` function joins observations that appear in both tables, based on a common key, **which need to be present in both tables**. It takes the following arguments, and returns an object of the same type as x.
 - x and y, a pair of tibbles or data frames to be joined.
 - by, which is a character vector of variables to join by. In cases where the key column name is different, a named vector can be used, for example, `by = c("key_x" = "key_y")`.

Primary key

ProductID	Type
PR-1	Computer
PR-2	Tablet
PR-3	Phone
PR-9	Headphones

Primary key

Foreign key

OrderID	Quantity	ProductID
OR-1	1	PR-1
OR-2	2	PR-2
OR-3	1	PR-3
OR-4	2	PR-4
OR-5	3	PR-1

```
i_j <- dplyr::inner_join(orders,products,by="ProductID")
i_j
#> # A tibble: 4 x 4
#>   OrderID  Quantity ProductID Type
#>   <chr>     <dbl> <chr>    <chr>
#> 1 OR-1       1 PR-1    Computer
#> 2 OR-2       2 PR-2    Tablet
#> 3 OR-3       1 PR-3    Phone
#> 4 OR-5       3 PR-1    Computer
```

`left_join(x,y)`

- A left join will keep all observations in the tibble x, even if there is no match in tibble y.
- This is a widely used function, given that it maintains all the observations in x.
- We can now show two examples based on the tibbles orders and products.

Primary key	
ProductID	Type
PR-1	Computer
PR-2	Tablet
PR-3	Phone
PR-9	Headphones

Primary key		Foreign key
OrderID	Quantity	ProductID
OR-1	1	PR-1
OR-2	2	PR-2
OR-3	1	PR-3
OR-4	2	PR-4
OR-5	3	PR-1

```

l_j1 <- dplyr::left_join(orders,products,by="ProductID")
l_j1
#> # A tibble: 5 x 4
#>   OrderID  Quantity ProductID Type
#>   <chr>     <dbl>  <chr>    <chr>
#> 1 OR-1       1 PR-1    Computer
#> 2 OR-2       2 PR-2    Tablet
#> 3 OR-3       1 PR-3    Phone
#> 4 OR-4       2 PR-4    <NA>
#> 5 OR-5       3 PR-1    Computer

```

Primary key

ProductID	Type
PR-1	Computer
PR-2	Tablet
PR-3	Phone
PR-9	Headphones

Primary key

Foreign key

OrderID	Quantity	ProductID
OR-1	1	PR-1
OR-2	2	PR-2
OR-3	1	PR-3
OR-4	2	PR-4
OR-5	3	PR-1

```
l_j3 <- dplyr::left_join(products, orders, by = "ProductID")
l_j3
#> # A tibble: 5 x 4
#>   ProductID Type     OrderID  Quantity
#>   <chr>      <chr>    <chr>     <dbl>
#> 1 PR-1       Computer OR-1        1
#> 2 PR-1       Computer OR-5        3
#> 3 PR-2       Tablet    OR-2        2
#> 4 PR-3       Phone     OR-3        1
#> 5 PR-9       Headphones <NA>       NA
```

Primary key	
ProductID	Type
PR-1	Computer
PR-2	Tablet
PR-3	Phone
PR-9	Headphones

Primary key		Foreign key
OrderID	Quantity	ProductID
OR-1	1	PR-1
OR-2	2	PR-2
OR-3	1	PR-3
OR-4	2	PR-4
OR-5	3	PR-1

```

l_j2 <- dplyr::left_join(orders,products,by="ProductID",keep=TRUE)
l_j2
#> # A tibble: 5 x 5
#>   OrderID  Quantity ProductID.x ProductID.y Type
#>   <chr>     <dbl> <chr>       <chr>      <chr>
#> 1 OR-1        1 PR-1        PR-1       Computer
#> 2 OR-2        2 PR-2        PR-2       Tablet
#> 3 OR-3        1 PR-3        PR-3       Phone
#> 4 OR-4        2 PR-4        <NA>       <NA>
#> 5 OR-5        3 PR-1        PR-1       Computer

```

`full_join()`

A full join keeps all observations in both x and y.

*Primary
key*

ProductID	Type
PR-1	Computer
PR-2	Tablet
PR-3	Phone
PR-9	Headphones

*Primary
key*

*Foreign
key*

OrderID	Quantity	ProductID
OR-1	1	PR-1
OR-2	2	PR-2
OR-3	1	PR-3
OR-4	2	PR-4
OR-5	3	PR-1

```
f_j1 <- dplyr::full_join(products,orders,by="ProductID")
f_j1
#> # A tibble: 6 x 4
#>   ProductID Type     OrderID  Quantity
#>   <chr>      <chr>    <chr>     <dbl>
#> 1 PR-1       Computer OR-1        1
#> 2 PR-1       Computer OR-5        3
#> 3 PR-2       Tablet    OR-2        2
#> 4 PR-3       Phone     OR-3        1
#> 5 PR-9       Headphones <NA>       NA
#> 6 PR-4       <NA>      OR-4        2
```

3. Filtering joins

- There are another set of joins that can be used to filter observations in tibble x, based on their relationship with values in another table.
- As the name suggests, these are filtering joins, and perform a similar task to the regular filter() function, except that information from two tables is used.

`semi_join(x,y)`

- This function will keep all the observations in x that have a matching column in y.

Primary key	
ProductID	Type
PR-1	Computer
PR-2	Tablet
PR-3	Phone
PR-9	Headphones

Primary key		Foreign key
OrderID	Quantity	ProductID
OR-1	1	PR-1
OR-2	2	PR-2
OR-3	1	PR-3
OR-4	2	PR-4
OR-5	3	PR-1

```
s_j1 <- dplyr::semi_join(orders, products, by="ProductID")
s_j1
#> # A tibble: 4 x 3
#>   OrderID  Quantity ProductID
#>   <chr>     <dbl>  <chr>
#> 1 OR-1       1 PR-1
#> 2 OR-2       2 PR-2
#> 3 OR-3       1 PR-3
#> 4 OR-5       3 PR-1
```

`anti_join(x,y)`

- This filtering function will keep all the observations in x that do not have a matching column in y.

Primary key	
ProductID	Type
PR-1	Computer
PR-2	Tablet
PR-3	Phone
PR-9	Headphones

Primary key		Foreign key
OrderID	Quantity	ProductID
OR-1	1	PR-1
OR-2	2	PR-2
OR-3	1	PR-3
OR-4	2	PR-4
OR-5	3	PR-1

```
a_j1 <- dplyr::anti_join(orders,products,by="ProductID")
a_j1
#> # A tibble: 1 × 3
#>   OrderID  Quantity ProductID
#>   <chr>     <dbl>  <chr>
#> 1 OR-4         2  PR-4
```

Mini-case: joining weather and energy data

- The aim of this example is to show how to combine related tibbles using `dplyr`, in order to explore new relationships.
- We use existing data from the package `aimsir17`, and in particular two tibbles:
 - `observations`, which contains hourly information from 2017 for 25 weather stations across Ireland, including the rainfall, temperature, humidity, mean sea level atmospheric pressure, wind speed, and wind direction.
 - `eirgrid17`, which holds information, recorded four times per hour, on energy demand and supply from the Irish Grid in 2017, including overall demand, energy generated, and wind energy generated.

eirgrid17

year	month	day	hour	minute	date	NIGeneration	NIDemand	NIWindAvailability	NIWindGeneration	IEGeneration	IEDemand	IEWindAvailability	IEWindGeneration
2017	1	1	0	0	2017-01-01 00:00:00	889.005	775.931	175.065	198.202	3288.57	2921.44	1064.79	1044.72
2017	1	1	0	15	2017-01-01 00:15:00	922.234	770.233	182.866	207.765	3282.12	2884.19	965.60	957.74
2017	1	1	0	30	2017-01-01 00:30:00	908.122	761.186	169.796	193.103	3224.27	2806.38	915.35	900.46
2017	1	1	0	45	2017-01-01 00:45:00	918.802	742.718	167.501	190.757	3171.27	2718.77	895.38	870.81
2017	1	1	1	0	2017-01-01 01:00:00	882.441	749.238	174.094	195.790	3190.28	2682.91	1028.03	998.31
2017	1	1	1	15	2017-01-01 01:15:00	848.863	742.455	189.922	212.956	3184.67	2649.87	1144.17	1119.12
2017	1	1	1	30	2017-01-01 01:30:00	842.778	726.472	222.139	244.569	3100.66	2578.31	1080.40	1056.76
2017	1	1	1	45	2017-01-01 01:45:00	808.910	709.353	233.321	245.045	3125.78	2555.87	1184.10	1165.61
2017	1	1	2	0	2017-01-01 02:00:00	797.183	697.191	281.727	298.625	3106.37	2498.73	1300.97	1275.92
2017	1	1	2	15	2017-01-01 02:15:00	754.976	684.101	259.120	285.191	3078.77	2441.95	1362.96	1337.95
2017	1	1	2	30	2017-01-01 02:30:00	795.311	668.003	272.766	306.300	2984.55	2400.41	1293.22	1266.24
2017	1	1	2	45	2017-01-01 02:45:00	780.488	651.126	261.728	290.701	2947.35	2365.22	1283.57	1257.01
2017	1	1	3	0	2017-01-01 03:00:00	777.013	640.060	286.211	308.586	2919.70	2358.54	1338.51	1314.98
2017	1	1	3	15	2017-01-01 03:15:00	780.512	628.341	266.933	293.743	2850.68	2306.28	1278.03	1253.87

observations

station	year	month	day	hour	date	rain	temp	rhum	msl	wdsp	wddir
ATHENRY	2017	1	1	0	2017-01-01 00:00:00	0.0	5.2	89	1021.9	8	320
ATHENRY	2017	1	1	1	2017-01-01 01:00:00	0.0	4.7	89	1022.0	9	320
ATHENRY	2017	1	1	2	2017-01-01 02:00:00	0.0	4.2	90	1022.1	8	320
ATHENRY	2017	1	1	3	2017-01-01 03:00:00	0.1	3.5	87	1022.5	9	330
ATHENRY	2017	1	1	4	2017-01-01 04:00:00	0.1	3.2	89	1022.7	8	330
ATHENRY	2017	1	1	5	2017-01-01 05:00:00	0.0	2.1	91	1023.3	8	330
ATHENRY	2017	1	1	6	2017-01-01 06:00:00	0.0	2.0	89	1023.5	7	330
ATHENRY	2017	1	1	7	2017-01-01 07:00:00	0.0	1.7	89	1024.4	7	340
ATHENRY	2017	1	1	8	2017-01-01 08:00:00	0.0	1.0	91	1025.0	7	330
ATHENRY	2017	1	1	9	2017-01-01 09:00:00	0.0	1.1	91	1026.1	8	330
ATHENRY	2017	1	1	10	2017-01-01 10:00:00	0.0	3.0	84	1026.8	9	320
ATHENRY	2017	1	1	11	2017-01-01 11:00:00	0.0	4.3	78	1027.0	12	350
ATHENRY	2017	1	1	12	2017-01-01 12:00:00	0.0	5.1	75	1027.4	11	360
ATHENRY	2017	1	1	13	2017-01-01 13:00:00	0.0	5.5	72	1027.4	12	360
ATHENRY	2017	1	1	14	2017-01-01 14:00:00	0.0	5.9	72	1027.6	11	360

aimsir17::eirgrid17

year
month
day
hour
minute
...
IEWindGeneration

eirgrid17_h

year
month
day
hour
IEWindGenerationH

summarize()

aimsir17::observations

year
month
day
hour
station
...
wdsp

left_join()

weather_energy

year
month
day
hour
IEWindGenerationH
station
wdsp

Create hourly summaries...

```
eirgrid17_h <- eirgrid17 %>%  
  dplyr::group_by(year,month,day,hour) %>%  
  dplyr::summarize(IEWindGenerationH=  
    mean(IEWindGeneration,  
         na.rm=T)) %>%  
  dplyr::ungroup()  
dplyr::slice(eirgrid17_h,1:4)  
#> # A tibble: 4 x 5  
#>   year month   day hour IEWindGenerationH  
#>   <dbl> <dbl> <int> <int>          <dbl>  
#> 1  2017     1     1     0        943.  
#> 2  2017     1     1     1       1085.  
#> 3  2017     1     1     2       1284.  
#> 4  2017     1     1     3       1254.
```

Use `left_join(x,y)`

```
obs1 <- observations %>%  
  dplyr::filter(!is.na(wdsp))  
  
weather_energy <- dplyr::left_join(eirgrid17_h,  
                                    obs1,  
                                    by=c("year",  
                                         "month",  
                                         "day",  
                                         "hour")) %>%  
  dplyr::select(year,month,day,hour,  
                IEWindGenerationH,  
                station,  
                wdsp)
```

Joined data

```
weather_energy
```


```
#> # A tibble: 201,430 x 7
#>   year month   day hour IEWindGenerationH station      wdsp
#>   <dbl> <dbl> <int> <int>          <dbl> <chr>      <dbl>
#> 1 2017     1     1     0 943. ATHENRY            8
#> 2 2017     1     1     0 943. BALLYHAISE         5
#> 3 2017     1     1     0 943. BELMULLET        13
#> 4 2017     1     1     0 943. CASEMENT           8
#> 5 2017     1     1     0 943. CLAREMORRIS        8
#> 6 2017     1     1     0 943. CORK AIRPORT       11
#> 7 2017     1     1     0 943. DUBLIN AIRPORT      12
#> 8 2017     1     1     0 943. DUNSANY             6
#> 9 2017     1     1     0 943. FINNER              12
#> 10 2017    1     1     0 943. GURTEEN             7
#> # ... with 201,420 more rows
```

Extract sample data

```
set.seed(100)
obs_sample <- weather_energy %>%
  dplyr::filter(station %in% c("MACE HEAD")) %>%
  dplyr::sample_n(300)

ggplot(obs_sample,aes(x=wdsp,y=IEWindGenerationH))+
  geom_point()+
  geom_jitter()+
  geom_smooth(method="lm")
```

Visualise Plot




Comparing Stations

```
corr_sum <- weather_energy %>%
  dplyr::group_by(station) %>%
  dplyr::summarize(Correlation=cor(wdsp,
  IEWindGenerationH)) %>%
  dplyr::arrange(desc(Correlation))

corr_sum
#> # A tibble: 23 x 2
#>   station      Correlation
#>   <chr>          <dbl>
#> 1 GURTEEN        0.806
#> 2 KNOCK AIRPORT  0.802
#> 3 MACE HEAD      0.800
#> 4 VALENTIA OBSERVATORY  0.778
#> 5 SHANNON AIRPORT  0.778
#> 6 MULLINGAR      0.771
#> 7 SherkinIsland    0.768
#> 8 CORK AIRPORT     0.768
#> 9 ROCHE'S POINT    0.767
#> 10 ATHENRY        0.764
#> # ... with 13 more rows
```

Showing correlations



Challenge

1. Based on the package `nycflights13`, which can be downloaded from CRAN, generate the following tibble based on the first three records from the tibble `flights`, and the airline name from `airlines`.

```
first_3a
#> # A tibble: 3 x 5
#>   time_hour          origin dest carrier name
#>   <dttm>            <chr>  <chr> <chr>   <chr>
#> 1 2013-01-01 05:00:00 EWR    IAH    UA      United Air Lines Inc.
#> 2 2013-01-01 05:00:00 LGA    IAH    UA      United Air Lines Inc.
#> 3 2013-01-01 05:00:00 JFK    MIA    AA      American Airlines Inc.
```