

Data Science for Operational Researchers Using R Online

8. Tibble Manipulation with `purrr` and `tidyverse`

Prof. Jim Duggan,
School of Computer Science
University of Galway.

https://github.com/JimDuggan/explore_or

R has numerous ways to iterate over elements of a list (or vector), and Hadley Wickham aimed to improve on and standardise that experience with the `purrr` package.

— Jared P. Lander ([Lander, 2017](#))

Overview



1. Overview

- A benefit of using the `tidyverse` is having the facility to combine tools from different packages, and switching between the use of lists and tibbles where appropriate.
- A common task is to divide a tibble into sub-groups, and perform operations on these.
- We have already seen how this can work using the package `dplyr`, which allows you to use the functions `group_by()` and `summarize()` to aggregate data
- The two functions we use with `purrr` are `dplyr:: group_split()` `tidyr::nest()`

2. group_split()

- This function, contained in the package dplyr, can be used to split a tibble into a list of tibbles, based on groupings specified by `group_by()`.
- This list can be processed using the `map()` family of functions.

Create a subset of mpg

```
set.seed(100)
test <- mpg %>%
  dplyr::select(manufacturer:displ,cty,class) %>%
  dplyr::filter(class %in% c("compact","midsize")) %>%
  dplyr::sample_n(5)

test
#> # A tibble: 5 x 5
#>   manufacturer model    displ     cty class
#>   <chr>        <chr>    <dbl>    <int> <chr>
#> 1 volkswagen   jetta     2.0      21 compact
#> 2 volkswagen   jetta     2.5      21 compact
#> 3 chevrolet    malibu    3.6      17 midsize
#> 4 volkswagen   gti       2.0      21 compact
#> 5 audi         a4        2.0      21 compact
```

Use group_by() and group_split()

```
test_s <- test %>%  
  dplyr::group_by(class) %>%  
  dplyr::group_split()
```

```
#> [[1]]  
#> # A tibble: 4 x 5  
#>   manufacturer model displ  cty class  
#>   <chr>        <chr> <dbl> <int> <chr>  
#> 1 volkswagen   jetta     2     21 compact  
#> 2 volkswagen   jetta    2.5    21 compact  
#> 3 volkswagen   gti      2     21 compact  
#> 4 audi         a4       2     21 compact
```

```
#> [[2]]  
#> # A tibble: 1 x 5  
#>   manufacturer model   displ  cty class  
#>   <chr>        <chr> <dbl> <int> <chr>  
#> 1 chevrolet    malibu   3.6    17 midsize
```

Processing the list with purrr::map_int()

```
test_s %>% purrr::map_int(~nrow(.x))  
#> [1] 4 1
```

```
#> [[1]]  
#> # A tibble: 4 x 5  
#>   manufacturer model displ  cty class  
#>   <chr>        <chr> <dbl> <int> <chr>  
#> 1 volkswagen   jetta     2     21 compact  
#> 2 volkswagen   jetta    2.5    21 compact  
#> 3 volkswagen   gti      2     21 compact  
#> 4 audi         a4       2     21 compact
```

```
#> [[2]]  
#> # A tibble: 1 x 5  
#>   manufacturer model   displ  cty class  
#>   <chr>        <chr> <dbl> <int> <chr>  
#> 1 chevrolet    malibu   3.6    17 midsize
```

More detailed example...

- Our goal is to calculate the **correlation coefficient** between two variables: mean sea level pressure and average wind speed.
- We simplify the dataset to daily values, where we take (1) the maximum wind speed (`wdsp`) recorded and (2) the average mean sea level pressure (`msl`).
- Our first task is to use `dplyr` to generate a summary tibble, and we also exclude any cases that have missing values, by combining `complete.cases()` within `filter()`.
- Note that the function `complete.cases()` returns a logical vector indicating which rows are complete.
- The new tibble is stored in the variable `d_data`.

```
d_data <- observations %>%  
  dplyr::filter(complete.cases(observations)) %>%  
  dplyr::group_by(station,month,day) %>%  
  dplyr::summarize(MaxWdsp=max(wdsp,na.rm=TRUE),  
                    DailyAverageMSL=mean(msl,na.rm=TRUE)) %>%  
  dplyr::ungroup()  
  
d_data  
#> # A tibble: 8,394 x 5  
#>   station month   day MaxWdsp DailyAverageMSL  
#>   <chr>     <dbl> <int>    <dbl>           <dbl>  
#> 1 ATHENRY     1     1      12       1027.  
#> 2 ATHENRY     1     2       8       1035.  
#> 3 ATHENRY     1     3       6       1032.  
#> 4 ATHENRY     1     4       4       1030.  
#> 5 ATHENRY     1     5       9       1029.  
#> 6 ATHENRY     1     6       9       1028.  
#> 7 ATHENRY     1     7       6       1032.  
#> 8 ATHENRY     1     8       9       1029.  
#> 9 ATHENRY     1     9      16       1015.  
#> 10 ATHENRY    1    10      13       1013.  
#> # ... with 8,384 more rows
```

```
cor7 <- d_data %>%
  dplyr::group_by(station) %>%
  dplyr::group_split() %>%
  purrr::map_df(~{
    corr <- cor(.x$MaxWdsp,.x$DailyAverageMSL)
    tibble(Station=first(.x$station),
           CorrCoeff=corr)
  }) %>%
  dplyr::arrange(CorrCoeff) %>%
  dplyr::slice(1:7)
```

```
cor7
#> # A tibble: 7 x 2
#>   Station          CorrCoeff
#>   <chr>            <dbl>
#> 1 SherkinIsland   -0.589
#> 2 VALENTIA OBSERVATORY -0.579
#> 3 ROCHE'S POINT    -0.540
#> 4 MACE HEAD        -0.539
#> 5 MOORE PARK       -0.528
```

Using summarize()

```
cor7_b <- d_data %>%
  dplyr::group_by(station) %>%
  dplyr::summarize(CorrCoeff=cor(MaxWdsp,DailyAverageMSL)) %>%
  dplyr::arrange(CorrCoeff) %>%
  dplyr::slice(1:7)

cor7_b
#> # A tibble: 7 x 2
#>   station          CorrCoeff
#>   <chr>              <dbl>
#> 1 SherkinIsland     -0.589
#> 2 VALENTIA OBSERVATORY -0.579
#> 3 ROCHE'S POINT      -0.540
#> 4 MACE HEAD           -0.539
#> 5 MOORE PARK          -0.528
```

3. `nest()`

- The function `nest()`, which is part of the package `tidyverse`, can be used to create a list column within a `tibble` that contains a `tibble`.
- Nesting generates one row for each defined group, which is identified using the function `group_by()`.
- The second column is named `data`, and is a list, and each list element contains all of the `tibble`'s data for a particular group.

```
data_n <- d_data %>%  
  dplyr::group_by(station) %>%  
  tidyverse::nest()  
  
data_n %>% head()  
#> # A tibble: 6 x 2  
#> # Groups:   station [6]  
#>   station      data  
#>   <chr>        <list>  
#> 1 ATHENRY      <tibble [365 x 4]>  
#> 2 BALLYHAISE   <tibble [365 x 4]>  
#> 3 BELMULLET   <tibble [365 x 4]>  
#> 4 CASEMENT     <tibble [365 x 4]>  
#> 5 CLAREMORRIS  <tibble [365 x 4]>  
#> 6 CORK AIRPORT <tibble [365 x 4]>
```

```
data_n %>% head()  
#> # A tibble: 6 x 2  
#> # Groups: station [6]  
#>   station      data  
#>   <chr>       <list>  
#> 1 ATHENRY     <tibble [365 x 4]>  
#> 2 BALLYHAISE  <tibble [365 x 4]>  
#> 3 BELMULLET  <tibble [365 x 4]>  
#> 4 CASEMENT    <tibble [365 x 4]>  
#> 5 CLAREMORRIS <tibble [365 x 4]>  
#> 6 CORK AIRPORT <tibble [365 x 4]>
```

```
data_n %>%  
  dplyr::pull(data) %>%  
  dplyr::first()  
  
#> # A tibble: 365 x 4  
#>   month   day MaxWdsp DailyAverageMSL  
#>   <dbl> <int> <dbl> <dbl>  
#> 1     1     1     12    1027.  
#> 2     1     2      8    1035.  
#> 3     1     3      6    1032.  
#> 4     1     4      4    1030.  
#> 5     1     5      9    1029.  
#> 6     1     6      9    1028.  
#> 7     1     7      6    1032.  
#> 8     1     8      9    1029.  
#> 9     1     9     16    1015.  
#> 10    1    10     13    1013.  
#> # ... with 355 more rows
```

4. Run linear regression model.

```
data_n <- data_n %>%  
  dplyr::mutate(LM=map(data,  
    ~lm(MaxWdsp~DailyAverageMSL,  
        data=.)))  
  
data_n %>%  
  head()  
#> # A tibble: 6 x 3  
#> # Groups:   station [6]  
#>   station      data          LM  
#>   <chr>       <list>        <list>  
#> 1 ATHENRY     <tibble [365 x 4]> <lm>  
#> 2 BALLYHAISE  <tibble [365 x 4]> <lm>  
#> 3 BELMULLET  <tibble [365 x 4]> <lm>  
#> 4 CASEMENT    <tibble [365 x 4]> <lm>  
#> 5 CLAREMORRIS <tibble [365 x 4]> <lm>  
#> 6 CORK AIRPORT <tibble [365 x 4]> <lm>
```

```

data_n %>%
  dplyr::filter(station=="BELMULLET") %>%
  dplyr::pull(LM) %>%
  dplyr::first() %>%
  summary()
#>
#> Call:
#> lm(formula = MaxWdsp ~ DailyAverageMSL, data = .)
#>
#> Residuals:
#>    Min     1Q Median     3Q    Max
#> -14.021 -4.069 -0.516  3.958 17.962
#>
#> Coefficients:
#>                               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)                242.786    26.365   9.21 <2e-16 ***
#> DailyAverageMSL           -0.222     0.026  -8.53  4e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 5.8 on 363 degrees of freedom
#> Multiple R-squared:  0.167, Adjusted R-squared:  0.165
#> F-statistic: 72.8 on 1 and 363 DF,  p-value: 4.03e-16

```

```
data_n <- data_n %>%  
  dplyr::mutate(RSq=map_db1(LM,~summary(.x)$r.squared)) %>%  
  dplyr::arrange(desc(RSq))  
data_n <- data_n %>% head(n=3)  
data_n  
#> # A tibble: 3 x 4  
#> # Groups:   station [3]  
#>   station          data           LM        RSq  
#>   <chr>         <list>        <list> <dbl>  
#> 1 SherkinIsland <tibble [365 x 4]> <lm>    0.347  
#> 2 VALENTIA OBSERVATORY <tibble [365 x 4]> <lm>    0.335  
#> 3 ROCHE'S POINT  <tibble [365 x 4]> <lm>    0.291
```

5. Additional purr functions

pluck()

- The function `pluck()` provides a generalized form of the `[[` operator and provides the means to index data structures in a flexible way.
- The arguments include `.x`, which is a vector, and a list of accessors for indexing into the object, which can include an integer position or a string name.
- Here are some examples.

```
library(ggplot2)
library(repurrrsive)

# Use pluck() to access the second element of an atomic vector
mpg %>% dplyr::pull(class) %>% unique() %>% purrr::pluck(2)
#> [1] "midsize"

# Use pluck() to access the director in the first list location
sw_films %>% purrr::pluck(1,"director")
#> [1] "George Lucas"
```

walk()

- The function `walk(.x,.f)` is similar to `map`, except that it returns the input `.x` and calls the function `.f` to generate a side effect.
- The side effect, for example, could be displaying information onto the screen, and no output value needs to be returned.

```
l <- list(el1=20,el2=30,el3=40)
o <- purrr::walk(l,~cat("Creating a side effect...\n"))
#> Creating a side effect...
#> Creating a side effect...
#> Creating a side effect...
str(o)
#> List of 3
#> $ el1: num 20
#> $ el2: num 30
#> $ el3: num 40
```

keep()

- The function `keep(.x,.f)` takes in a list `.x` and, based on the evaluation of a predicate function, will either keep or discard list element.
- In effect, it provides a way to filter a list. Here, we can filter those movies that have George Lucas as a director, and then confirm the result using `walk()`.

```
o <- sw_films %>% keep(~.x$director=="George Lucas")
purrr::walk(o,~cat(.x$director," ==> Title =",.x$title,"\n"))
#> George Lucas ==> Title = A New Hope
#> George Lucas ==> Title = Attack of the Clones
#> George Lucas ==> Title = The Phantom Menace
#> George Lucas ==> Title = Revenge of the Sith
```

3. Generate the following daily summaries of rainfall and mean sea level pressure, for all the weather stations in `aimsir17::observations`, and only consider observations with no missing values.

```
#> `summarize()` has grouped output by 'station', 'month'. You can  
#> override using the `$.groups` argument.
```

```
d_sum  
#> # A tibble: 8,394 x 5  
#>   station month   day TotalRain AvrMSL  
#>   <chr>     <dbl> <int>      <dbl>    <dbl>  
#> 1 ATHENRY     1     1        0.2    1027.  
#> 2 ATHENRY     1     2        0       1035.  
#> 3 ATHENRY     1     3        0       1032.  
#> 4 ATHENRY     1     4        0       1030.  
#> 5 ATHENRY     1     5        0.1     1029.  
#> 6 ATHENRY     1     6        18      1028.  
#> 7 ATHENRY     1     7        1.4     1032.  
#> 8 ATHENRY     1     8        1.2     1029.  
#> 9 ATHENRY     1     9        5.4     1015.  
#> 10 ATHENRY    1    10        0.7    1013.  
#> # ... with 8,384 more rows
```

Next, using the tibble `d_sum` as input, generate the top 6 correlations between `TotalRain` and `AvrMSL` using `group_split()` and `map_df()`. Here are the results you should find.

```
cors
#> # A tibble: 6 × 2
#>   Station          Corr
#>   <chr>           <dbl>
#> 1 MOORE PARK     -0.496
#> 2 MULLINGAR      -0.496
#> 3 CLAREMORRIS    -0.484
#> 4 VALENTIA OBSERVATORY -0.464
#> 5 KNOCK AIRPORT   -0.459
#> 6 OAK PARK        -0.456
```