

R has numerous ways to iterate over elements of a list (or vector), and Hadley Wickham aimed to improve on and standardise that experience with the `purrr` package.

— Jared P. Lander (Lander, 2017)

Data Science for Operational Researchers using R

08 – `purrr`

https://github.com/JimDuggan/explore_or

Overview

- We earlier used the functional `lapply()` to iterate over an atomic vector, list, or data frame.
- The tidyverse package `purrr` provides a comprehensive set of functions that can be used to iterate over data structures
- It also integrates with other elements of the tidyverse, for example, the package `dplyr`.
- Topics
 - `map()` family of functions
 - Integrating with `dplyr` to process tibble
 - Mini-case with linear models.

(1) The `map()` family of functions

- The idea of the `map()` function is to provide a mechanism to iterate over an input list or a vector.
- It applies a function to each input element, and returns the result within a list that is exactly the same length as the input.
- We call these functions functionals, as they accept a function as an argument, and use that function in order to generate the output




FIGURE 10.1 The `map` function in `purrr`

General format of map(.x, .f)

- .x is a list, or an atomic vector.
- If .x has named elements, the return value will preserve those names.
- .f can be a function, formula, or vector.

```
library(purrr)

o1 <- purrr::map(c(1,2,3,2),function(x)x^2)
str(o1)
#> List of 4
#> $ : num 1
#> $ : num 4
#> $ : num 9
#> $ : num 4
```

Using a formula

- A formula (which is defined by ~) can be used as a function argument, and it is essentially a function shortcut in purrr, and commonly used in map()
- ~ marks the start of the function, and .x as the function parameter, where . can also be used to represent the function parameter.

```
o2 <- purrr::map(c(1,2,3,2),~.x^2)
str(o2)
#> List of 4
#> $ : num 1
#> $ : num 4
#> $ : num 9
#> $ : num 4

o3 <- purrr::map(c(1,2,3,2),~.^2)
str(o3)
#> List of 4
#> $ : num 1
#> $ : num 4
#> $ : num 9
#> $ : num 4
```

Using map() to select a list element

map() functions can be used to select a list element by name, and these are beneficial when working with lists that are deeply nested (Wickham, 2019).

```
library(repurrrsive)
dirs <- sw_films %>% purrr::map("director") %>% unique()
str(dirs)
#> List of 4
#> $ : chr "George Lucas"
#> $ : chr "Richard Marquand"
#> $ : chr "Irvin Kershner"
#> $ : chr "J. J. Abrams"
```

Additional map_* functions

While `map()` will always return a list, there may be circumstances where different formats are required, for example, an atomic vector. To address this, `purrr` provides a set of additional functions that specify the result type. These include:

- `map_dbl()`, which returns an atomic vector of type double.
- `map_chr()`, which returns an atomic vector of type character.
- `map_lgl()`, which returns an atomic vector of type logical.
- `map_int()`, which returns an atomic vector of type integer.
- `map_df()`, which returns a data frame or tibble.

map_dbl()

Here, we can process a number of columns from the data frame `mtcars` and return the average of each column as an atomic vector. Note that because a data frame is also a list, we can use it as an input to the `map` family of functions.

```
library(dplyr)
library(purrr)
mtcars %>%
  dplyr::select(mpg, cyl, disp) %>%
  purrr::map_dbl(mean)
#>      mpg      cyl    disp
#> 20.091  6.188 230.722
```


map_chr()

If we wanted to extract the director names from `sw_films` as an atomic vector, we can do this with `map_chr()`. Note that in this example, we use a formula to specify the function shortcut, and the argument is accessed using `.x`.

```
library(repurrrsive)
library(purrr)
sw_films %>%
  purrr::map_chr(~.x$director) %>%
  unique()
#> [1] "George Lucas"      "Richard Marquand" "Irvin Kershner"
#> [4] "J. J. Abrams"
```

map_lgl()

Here, we process a number of the columns in `mpg` to test whether the columns are numeric. Here, we use the anonymous function option.

```
library(ggplot2)
library(purrr)
library(dplyr)
mpg %>%
  dplyr::select(manufacturer:cyl) %>%
  purrr::map_lgl(function(x)is.numeric(x))
```

#> manufacturer	model	displ	year	cyl
#> FALSE	FALSE	TRUE	TRUE	TRUE

map_int()

In this example, we select a number of numeric columns from `mpg`, and then use `map_int()` to count the number of observations that are greater than the mean in each of the three columns. An atomic vector of integers is returned.

```
library(ggplot2)
library(dplyr)
library(purrr)

mpg %>%
  dplyr::select(displ, cty, hwy) %>%
  purrr::map_int(~sum(.x>mean(.x)))
#> displ    cty    hwy
#>    107    118    129
```

map_df()

The function `map_df()` creates a new data frame or tibble based on the input list. A tibble is specified within the function, and as `map_df()` iterates through the input list, rows will be added to the tibble with the specified values. In

```
library(repurrrsive)
library(purrr)
library(dplyr)

sw_films %>%
  purrr::map_df(~tibble(ID=.x$episode_id,
                        Title=.x$title,
                        Director=.x$director,
                        ReleaseDate=as.Date(.x$release_date))) %>%
  dplyr::arrange(ID)
```

```
#> # A tibble: 7 x 4
#>       ID Title                                Director      ReleaseDate
#>   <int> <chr>                                <chr>         <date>
#> 1     1 1 The Phantom Menace                George Lucas  1999-05-19
#> 2     2 2 Attack of the Clones                George Lucas  2002-05-16
#> 3     3 3 Revenge of the Sith                George Lucas  2005-05-19
#> 4     4 4 A New Hope                          George Lucas  1977-05-25
#> 5     5 5 The Empire Strikes Back            Irvin Kershner 1980-05-17
#> 6     6 6 Return of the Jedi                  Richard Marquand 1983-05-25
#> 7     7 7 The Force Awakens                  J. J. Abrams  2015-12-11
```

Iterating over 2 inputs – map2()

The function `map2()` allows for two inputs, and these are then represented as arguments by `.x` and `.y`.

```
means <- c(10,20,30)
sds    <- c(2,4,7)

purrr::map2(means,sds,~rnorm(5,.x,.y)) %>% str()
#> List of 3
#> $ : num [1:5] 10.42 6.78 10.54 12.05 7.1
#> $ : num [1:5] 28 21.6 10.4 31.1 18.2
#> $ : num [1:5] 32 35.3 30.6 25.5 28.3
```

Iterating over multiple inputs with pmap()

pmap() can take a list containing any number of arguments, and process these elements within the function using the symbols ..1, ..2 which represent the first, second, and additional arguments

```
params <- list(means = c(10,20,30),  
              sds   = c(2,4,7),  
              n     = c(4,5,6))
```

```
purrr::pmap(params,  
             ~rnorm(n      = ..3,  
                   mean   = ..1,  
                   sd     = ..2)) %>%  
  str()  
#> List of 3  
#> $ : num [1:4] 12.68 10.27 5.06 11.31  
#> $ : num [1:5] 20.1 18.5 21 14.9 20.2  
#> $ : num [1:6] 28 25 34.4 35.6 36.7 ...
```

pmap() example

```
set.seed(100)
grades <- tibble(ID=paste0("S-",10:15),
                 Subject1=rnorm(6,70,10),
                 Subject2=rnorm(6,60,20),
                 Subject3=rnorm(6,50,15))
```

grades

```
#> # A tibble: 6 x 4
```

```
#>   ID      Subject1 Subject2 Subject3
#>   <chr>    <dbl>    <dbl>    <dbl>
#> 1 S-10      65.0      48.4      47.0
#> 2 S-11      71.3      74.3      61.1
#> 3 S-12      69.2      43.5      51.9
#> 4 S-13      78.9      52.8      49.6
#> 5 S-14      71.2      61.8      44.2
#> 6 S-15      73.2      61.9      57.7
```


pmap() example 2

```
grades1 <- grades %>%  
  dplyr::mutate(Summary=pmap_chr(grades,  
                                ~paste0("ID=",  
                                         ..1,  
                                         " Max=",  
                                         round(max(..2,..3,..4),2))))
```

grades1

```
#> # A tibble: 6 x 5
```

```
#>   ID      Subject1 Subject2 Subject3 Summary
```

```
#>   <chr>      <dbl>      <dbl>      <dbl> <chr>
```

```
#> 1 S-10      65.0       48.4       47.0 ID=S-10 Max=64.98
```

```
#> 2 S-11      71.3       74.3       61.1 ID=S-11 Max=74.29
```

```
#> 3 S-12      69.2       43.5       51.9 ID=S-12 Max=69.21
```

```
#> 4 S-13      78.9       52.8       49.6 ID=S-13 Max=78.87
```

```
#> 5 S-14      71.2       61.8       44.2 ID=S-14 Max=71.17
```

```
#> 6 S-15      73.2       61.9       57.7 ID=S-15 Max=73.19
```


2. Integrating purrr with dplyr and tidyr

- A benefit of using the tidyverse is having the facility to combine tools from different packages, and switching between the use of lists and tibbles where appropriate.
- A common task is to divide a tibble into sub-groups, and perform operations on these
- We first define a test data set

```
set.seed(100)
test <- mpg %>%
  dplyr::select(manufacturer:displ,cty,class) %>%
  dplyr::filter(class %in% c("compact","midsize")) %>%
  dplyr::sample_n(5)

test
#> # A tibble: 5 x 5
#>   manufacturer model  displ   cty class
#>   <chr>         <chr> <dbl> <int> <chr>
#> 1 volkswagen   jetta     2     21 compact
#> 2 volkswagen   jetta    2.5     21 compact
#> 3 chevrolet    malibu    3.6     17 midsize
#> 4 volkswagen   gti       2     21 compact
#> 5 audi         a4        2     21 compact
```

group_split()

- Takes a tibble
- And grouping variable(s)
- Generates a list of tibbles

```
test_s %>% purrr::map_int(~nrow(.x))  
#> [1] 4 1
```

```
test_s <- test %>%  
  dplyr::group_by(class) %>%  
  dplyr::group_split()  
  
test_s  
#> <list_of<  
#>   tbl_df<  
#>     manufacturer: character  
#>     model       : character  
#>     displ       : double  
#>     cty         : integer  
#>     class       : character  
#>   >  
#> >[2]>  
#> [[1]]  
#> # A tibble: 4 x 5  
#>   manufacturer model displ  cty class  
#>   <chr>         <chr> <dbl> <int> <chr>  
#> 1 volkswagen   jetta    2     21 compact  
#> 2 volkswagen   jetta    2.5   21 compact  
#> 3 volkswagen   gti      2     21 compact  
#> 4 audi         a4       2     21 compact  
#>  
#> [[2]]  
#> # A tibble: 1 x 5  
#>   manufacturer model displ  cty class  
#>   <chr>         <chr> <dbl> <int> <chr>  
#> 1 chevrolet    malibu   3.6   17  midsize
```

Exploring data

- Our goal is to calculate the correlation coefficient between two variables at each station: mean sea level pressure and average wind speed.
- We simplify the dataset to **daily values**, where we take (1) the maximum wind speed (wdsp) recorded and (2) the average mean sea level pressure (msl).
- Our first task is to use dplyr to generate a summary tibble, and we also exclude any cases that have missing values, by combining `complete.cases()` within filter

```
d_data <- observations %>%
  dplyr::filter(complete.cases(observations)) %>%
  dplyr::group_by(station, month, day) %>%
  dplyr::summarize(MaxWdsp=max(wdsp, na.rm=TRUE),
                  DailyAverageMSL=mean(msl, na.rm=TRUE)) %>%
  dplyr::ungroup()

d_data
#> # A tibble: 8,394 x 5
#>   station month   day MaxWdsp DailyAverageMSL
#>   <chr>    <dbl> <int>    <dbl>         <dbl>
#> 1 ATHENRY     1     1      12         1027.
#> 2 ATHENRY     1     2       8         1035.
#> 3 ATHENRY     1     3       6         1032.
#> 4 ATHENRY     1     4       4         1030.
#> 5 ATHENRY     1     5       9         1029.
#> 6 ATHENRY     1     6       9         1028.
#> 7 ATHENRY     1     7       6         1032.
#> 8 ATHENRY     1     8       9         1029.
#> 9 ATHENRY     1     9      16         1015.
#> 10 ATHENRY    1    10      13         1013.
#> # ... with 8,384 more rows
```

Process the results

```
cor7 <- d_data %>%
  dplyr::group_by(station) %>%
  dplyr::group_split() %>%
  purrr::map_df(~{
    corr <- cor(.x$MaxWdsp,.x$DailyAverageMSL)
    tibble(Station=first(.x$station),
           CorrCoeff=corr)
  }) %>%
  dplyr::arrange(CorrCoeff) %>%
  dplyr::slice(1:7)
```

```
cor7
#> # A tibble: 7 x 2
#>   Station                CorrCoeff
#>   <chr>                  <dbl>
#> 1 SherkinIsland          -0.589
#> 2 VALENTIA OBSERVATORY   -0.579
#> 3 ROCHES POINT           -0.540
#> 4 MACE HEAD              -0.539
#> 5 MOORE PARK             -0.528
#> 6 SHANNON AIRPORT        -0.524
#> 7 CORK AIRPORT           -0.522
```

Alternative solution...

```
cor7_b <- d_data %>%  
  dplyr::group_by(station) %>%  
  dplyr::summarize(CorrCoeff=cor(MaxWdsp,DailyAverageMSL)) %>%  
  dplyr::arrange(CorrCoeff) %>%  
  dplyr::slice(1:7)
```

```
cor7_b  
#> # A tibble: 7 x 2  
#>   station      CorrCoeff  
#>   <chr>         <dbl>  
#> 1 SherkinIsland -0.589  
#> 2 VALENTIA OBSERVATORY -0.579  
#> 3 ROCHES POINT -0.540  
#> 4 MACE HEAD -0.539  
#> 5 MOORE PARK -0.528  
#> 6 SHANNON AIRPORT -0.524  
#> 7 CORK AIRPORT -0.522
```

nest()

- `nest()`, which is part of the package `tidyr`, can be used to create a list column within a tibble that contains a data frame.
- Nesting generates one row for each defined group, which is identified using the function `group_by()`.
- The second column is named `data`, and is a list, and each list element contains all of the tibble's data for a particular group.

```
d_data
#> # A tibble: 8,394 x 5
#>   station month   day MaxWdsp DailyAverageMSL
#>   <chr>    <dbl> <int>   <dbl>         <dbl>
#> 1 ATHENRY      1     1     12         1027.
#> 2 ATHENRY      1     2      8         1035.
```

```
data_n <- d_data %>%
  dplyr::group_by(station) %>%
  tidyr::nest()
```

```
data_n %>% head()
#> # A tibble: 6 x 2
#> # Groups:   station [6]
#>   station      data
#>   <chr>      <list>
#> 1 ATHENRY    <tibble [365 x 4]>
#> 2 BALLYHAISE <tibble [365 x 4]>
#> 3 BELMULLET  <tibble [365 x 4]>
#> 4 CASEMENT   <tibble [365 x 4]>
#> 5 CLAREMORRIS <tibble [365 x 4]>
#> 6 CORK AIRPORT <tibble [365 x 4]>
```


The **data** column

- Here, the tibble `data_n` contains two columns, with a row for each weather station (the first six rows are shown here).
- All of the data for each weather station is stored in the respective cell in the column **data**.
- `first()` is a wrapper around the list operator `[[` that returns the first value in a list.
- Note the group name does not appear in `data` (it is already in the first column)

```
data_n %>% head()
#> # A tibble: 6 x 2
#> # Groups:   station [6]
#>   station      data
#>   <chr>      <list>
#> 1 ATHENRY    <tibble [365 x 4]>
```

```
data_n %>%
  dplyr::pull(data) %>%
  dplyr::first()
#> # A tibble: 365 x 4
#>   month   day MaxWdsp DailyAverageMSL
#>   <dbl> <int>   <dbl>         <dbl>
#> 1     1     1     12         1027.
#> 2     1     2      8         1035.
#> 3     1     3      6         1032.
#> 4     1     4      4         1030.
#> 5     1     5      9         1029.
#> 6     1     6      9         1028.
#> 7     1     7      6         1032.
#> 8     1     8      9         1029.
#> 9     1     9     16         1015.
#> 10    1    10     13         1013.
#> # ... with 355 more rows
```

Generating a linear model for each row using map()

```
data_n <- data_n %>%
  dplyr::mutate(LM=map(data,
                      ~lm(MaxWdsp~DailyAverageMSL,
                          data=.)))

data_n %>%
  head()
#> # A tibble: 6 x 3
#> # Groups:   station [6]
#>   station      data      LM
#>   <chr>      <list>    <list>
#> 1 ATHENRY   <tibble [365 x 4]> <lm>
#> 2 BALLYHAISE <tibble [365 x 4]> <lm>
#> 3 BELMULLET <tibble [365 x 4]> <lm>
#> 4 CASEMENT  <tibble [365 x 4]> <lm>
#> 5 CLAREMORRIS <tibble [365 x 4]> <lm>
#> 6 CORK AIRPORT <tibble [365 x 4]> <lm>
```


Exploring the result for “BELMULLET”

```
data_n %>%
  dplyr::filter(station=="BELMULLET") %>%
  dplyr::pull(LM) %>%
  dplyr::first() %>%
  summary()
#>
#> Call:
#> lm(formula = MaxWdsp ~ DailyAverageMSL, data = .)
#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
#> -14.021  -4.069  -0.516   3.958  17.962
#>
#> Coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    242.786     26.365   9.21  <2e-16 ***
#> DailyAverageMSL  -0.222      0.026  -8.53   4e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 5.8 on 363 degrees of freedom
#> Multiple R-squared:  0.167, Adjusted R-squared:  0.165
#> F-statistic: 72.8 on 1 and 363 DF, p-value: 4.03e-16
```

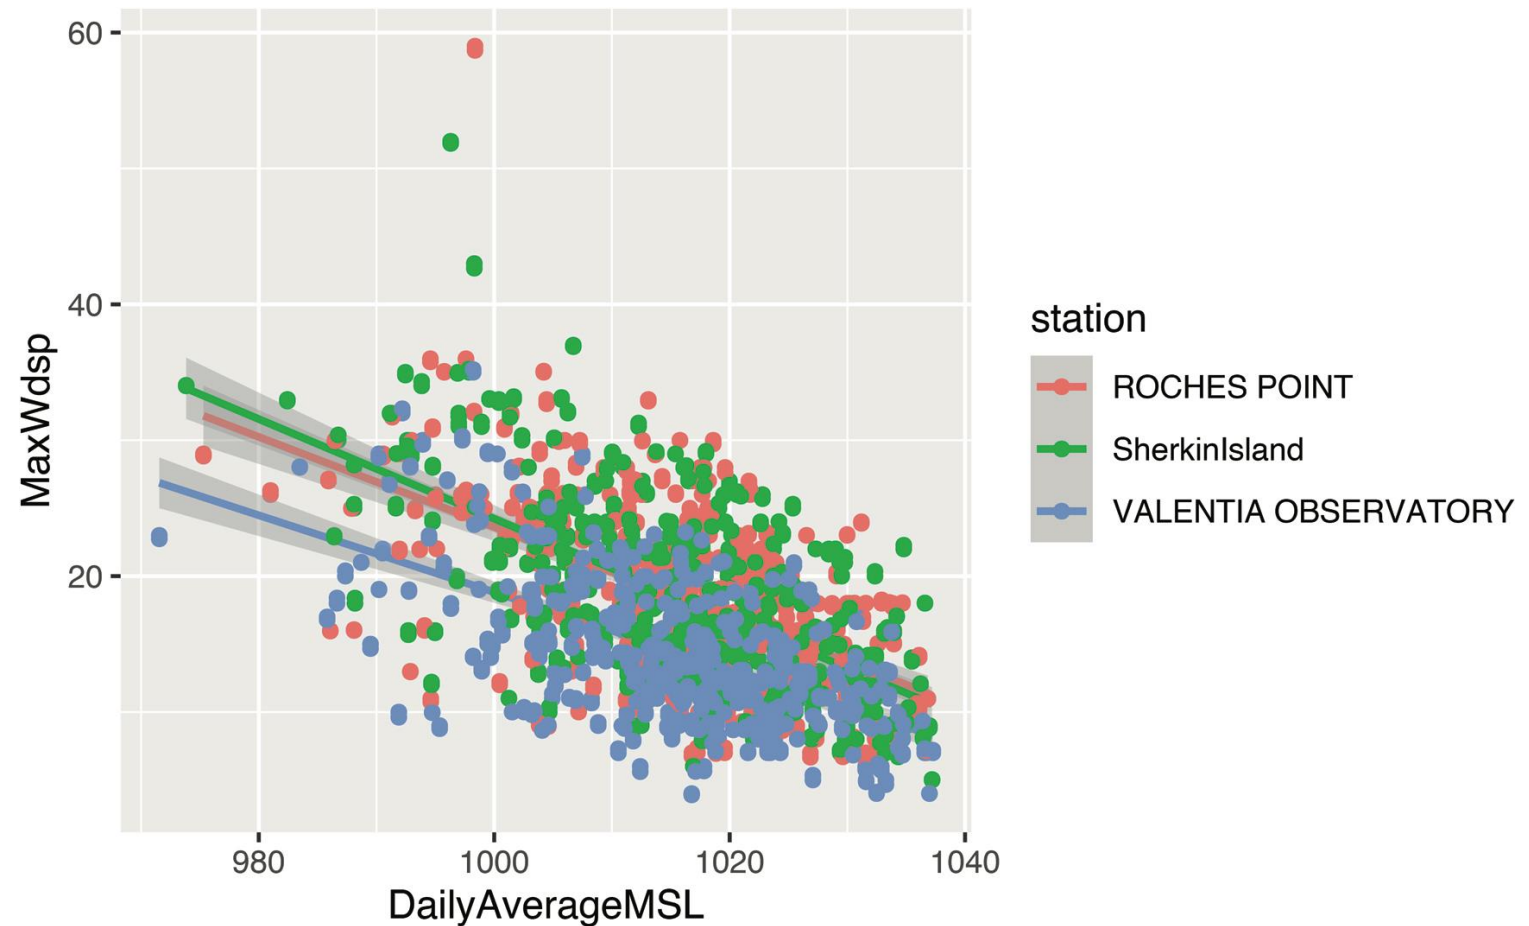
Note, probably not
needed...

Add R² value as a column

```
data_n <- data_n %>%
  dplyr::mutate(RSq=map_dbl(LM,~summary(.x)$r.squared)) %>%
  dplyr::arrange(desc(RSq))
data_n <- data_n %>% head(n=3)
data_n
#> # A tibble: 3 x 4
#> # Groups:   station [3]
#>   station          data          LM      RSq
#>   <chr>          <list>        <list> <dbl>
#> 1 SherkinIsland <tibble [365 x 4]> <lm>    0.347
#> 2 VALENTIA OBSERVATORY <tibble [365 x 4]> <lm>    0.335
#> 3 ROCHES POINT <tibble [365 x 4]> <lm>    0.291
```

```
data1 <- d_data %>%
  dplyr::filter(station %in% dplyr::pull(data_n,station))

ggplot(data1,aes(x=DailyAverageMSL,y=MaxWdsp,color=station))+
  geom_point()+geom_smooth(method="lm")+geom_jitter()
```



pluck()

- The function `pluck()` provides a generalized form of the `[[` operator and provides the means to index data structures in a flexible way.
- The arguments include `.x`, which is a vector, and a list of accessors for indexing into the object, which can include an integer position or a string name.

```
library(ggplot2)
library(repurrrsive)

# Use pluck() to access the second element of an atomic vector
mpg %>% dplyr::pull(class) %>% unique() %>% purrr::pluck(2)
#> [1] "midsize"

# Use pluck() to access the director in the first list location
sw_films %>% purrr::pluck(1,"director")
#> [1] "George Lucas"
```

walk()

- `walk(.x,.f)` is similar to `map`, except that it returns the input `.x` and calls the function `.f` to generate a side effect.
- The side effect, for example, could be displaying information onto the screen.

```
l <- list(el1=20,el2=30,el3=40)
o <- purrr::walk(l,~cat("Creating a side effect...\n"))
#> Creating a side effect...
#> Creating a side effect...
#> Creating a side effect...
str(o)
#> List of 3
#> $ el1: num 20
#> $ el2: num 30
#> $ el3: num 40
```

keep()

- The function `keep(.x,.f)` takes in a list `.x` and, based on the evaluation of a predicate function, will either keep or discard list element.
- In effect, it provides a way to filter a list.

```
o <- sw_films %>% keep(~.x$director=="George Lucas")
purrr::walk(o,~cat(.x$director," ==> Title =",.x$title,"\n"))
#> George Lucas ==> Title = A New Hope
#> George Lucas ==> Title = Attack of the Clones
#> George Lucas ==> Title = The Phantom Menace
#> George Lucas ==> Title = Revenge of the Sith
```

invoke_map()

- Provides a feature to call a list of functions with a list of parameters.
- It is a wrapper around the R function `do.call()`.
- The function takes a list of functions (where the function names are contained in a character vector), and an argument list, where the names of the function arguments are contained in the list.
- Function is now deprecated (see advice)

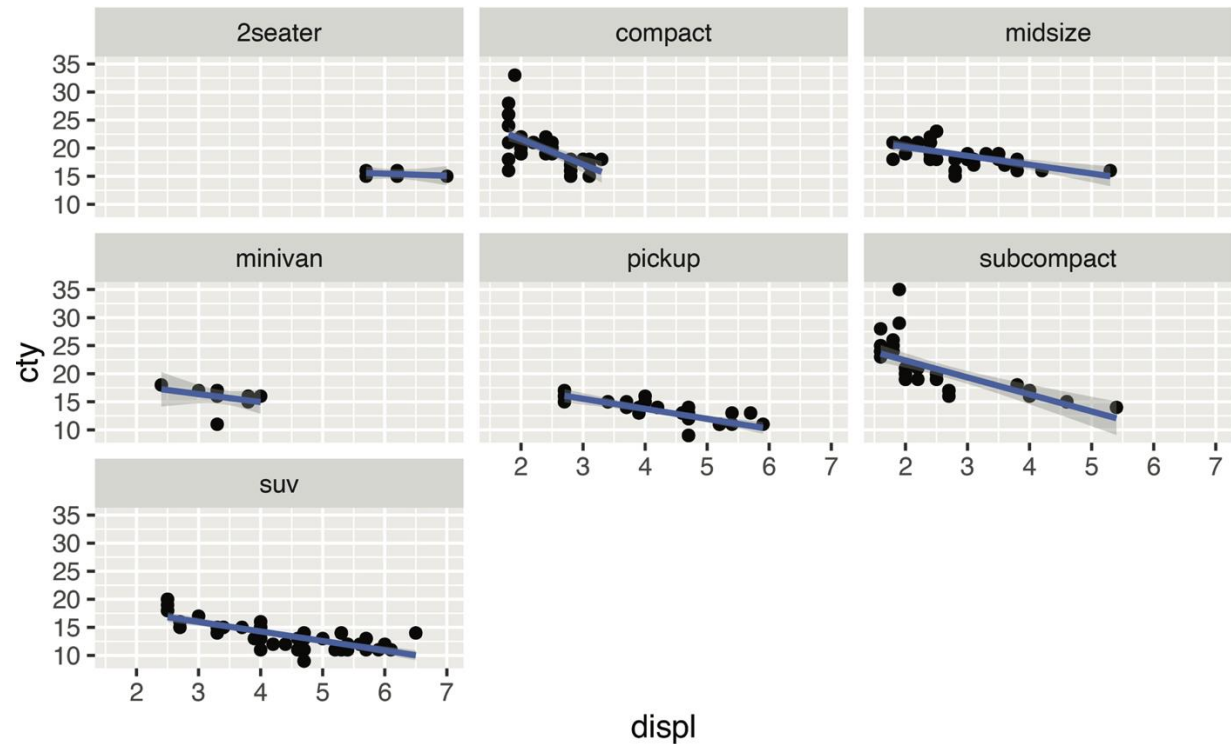
```
f <- c("min","max","sum")
l <- list(
  arg1=list(x=1:3),
  arg2=list(x=10:12),
  arg3=list(x=1:4)
)

str(purrr::invoke_map(f,l))
#> Warning: `invoke_map()` was deprecated in purrr 1.0.0.
#> i Please use map() + exec() instead.
#> This warning is displayed once every 8 hours.
#> Call `lifecycle::last_lifecycle_warnings()` to see where this
#> warning was generated.
#> List of 3
#> $ : int 1
#> $ : int 12
#> $ : int 10
```

Mini-case - Linear models

- Here, we generate a collection of linear models from the mpg tibble, and explore the relationship between a dependent variable (cty) and a dependent variable (displ).
- In this example, we will partition the data by vehicle class, and so build seven different linear models.
- Before starting the modelling workflow, we can visualize the different linear models using ggplot().

```
ggplot(mpg, aes(x=displ, y=cty)) +  
  geom_point() + geom_smooth(method="lm") +  
  facet_wrap(~class)
```



Exercise 1

1. Create the following tibble with the three columns shown, using the functions `keep()` and `map_df()`, in order to provide a tabular view of the list `repurrrsive::sw_vehicles`. Note that possible invalid values of `length` in `sw_vehicles` include “unknown”, and any of these should be removed prior to creating the data frame.

```
#> # A tibble: 6 x 3
#>   Name                Model                Length
#>   <chr>              <chr>              <dbl>
#> 1 C-9979 landing craft C-9979 landing craft      210
#> 2 SPHA               Self-Propelled Heavy Artillery 140
#> 3 Clone turbo tank    HAVw A6 Juggernaut        49.4
#> 4 Sand Crawler        Digger Crawler            36.8
#> 5 Multi-Troop Transport Multi-Troop Transport      31
#> 6 Sail barge          Modified Luxury Sail Barge   30
```

Overall approach

- Create a nested tibble (grouped by class), and this will have two columns, class and data, where the data column is created by the call to `nest()`.
- Using `mutate()`, add a new column that will store the results of each linear model.
- Process the linear model results to extract a measure of model performance (r-squared) to see how well each regression model explains observed data.
- Arrange the tibble so that the best model scores are shown first.
- Add a column that will store the plots for each model.

Use nest() to reconfigure tibble

```
mpg1 <- mpg %>%  
  dplyr::group_by(class) %>%  
  tidyr::nest()
```

```
mpg1
```

```
#> # A tibble: 7 x 2  
#> # Groups:   class [7]  
#>   class      data  
#>   <chr>    <list>  
#> 1 compact <tibble [47 x 10]>  
#> 2 midsize <tibble [41 x 10]>  
#> 3 suv     <tibble [62 x 10]>  
#> 4 2seater <tibble [5 x 10]>  
#> 5 minivan <tibble [11 x 10]>  
#> 6 pickup  <tibble [33 x 10]>  
#> 7 subcompact <tibble [35 x 10]>
```

Perform regression on each vehicle class

```
mpg1 <- mpg1 %>%  
  dplyr::mutate(LM=map(data,~lm(cty~displ,data=.x)))  
mpg1  
#> # A tibble: 7 x 3  
#> # Groups:   class [7]  
#>   class      data      LM  
#>   <chr>    <list>    <list>  
#> 1 compact <tibble [47 x 10]> <lm>  
#> 2 midsize <tibble [41 x 10]> <lm>  
#> 3 suv     <tibble [62 x 10]> <lm>  
#> 4 2seater <tibble [5 x 10]>  <lm>  
#> 5 minivan <tibble [11 x 10]> <lm>  
#> 6 pickup  <tibble [33 x 10]> <lm>  
#> 7 subcompact <tibble [35 x 10]> <lm>
```

Show one of the model results

```
mpg1 %>%
  dplyr::filter(class=="suv") %>% # Select row where class == suv"
  dplyr::pull(LM) %>%             # Get the column "LM"
  dplyr::first() %>%             # Extract the lm object
  summary()                       # Call the summary function
#>
#> Call:
#> lm(formula = cty ~ displ, data = .x)
#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
#> -4.087 -1.027 -0.087  1.096  3.967
#>
#> Coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   21.060      0.893    23.6  < 2e-16 ***
#> displ        -1.696      0.195    -8.7  3.2e-12 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 1.62 on 60 degrees of freedom
#> Multiple R-squared:  0.558, Adjusted R-squared:  0.55
#> F-statistic: 75.7 on 1 and 60 DF, p-value: 3.17e-12
```

Extract R² value into new column

```
mpg1 <- mpg1 %>%  
  dplyr::mutate(RSquared=map_dbl(LM,~summary(.x)$r.squared)) %>%  
  dplyr::arrange(desc(RSquared))  
  
mpg1  
#> # A tibble: 7 x 4  
#> # Groups:   class [7]  
#>   class      data          LM      RSquared  
#>   <chr>      <list>        <list>    <dbl>  
#> 1 suv        <tibble [62 x 10]> <lm>      0.558  
#> 2 subcompact <tibble [35 x 10]> <lm>      0.527  
#> 3 pickup     <tibble [33 x 10]> <lm>      0.525  
#> 4 compact    <tibble [47 x 10]> <lm>      0.358  
#> 5 midsize    <tibble [41 x 10]> <lm>      0.339  
#> 6 2seater    <tibble [5 x 10]>  <lm>      0.130  
#> 7 minivan    <tibble [11 x 10]> <lm>      0.124
```

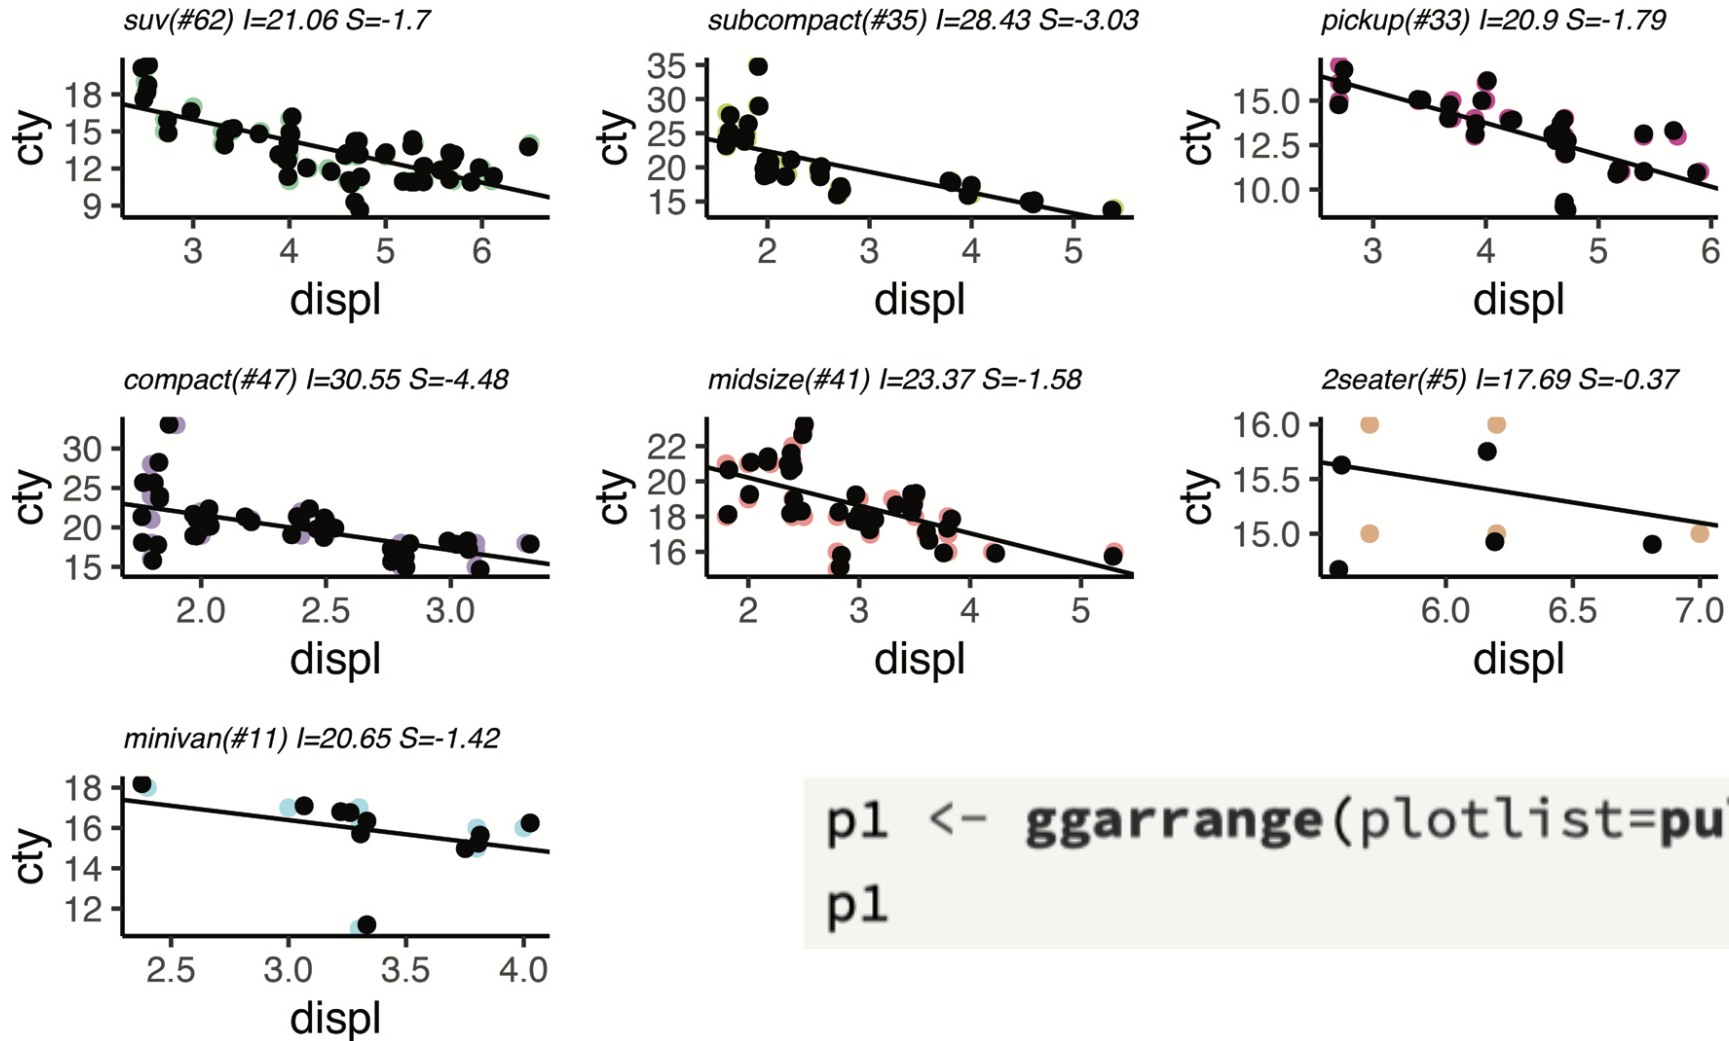
Create custom plots via `map2()`

```
library(randomcolorR) # Functions to generate different colors
set.seed(100) # For random color generation

mpg1 <- mpg1 %>%
  dplyr::mutate(Plots=map2(class,data, ~{
    # (1) run linear model
    m <- lm(cty~displ,data=.y)
    # (2) Extract coefficients
    intercept <- round(coef(m)[1],2)
    slope <- round(coef(m)[2],2)
    # (3) return the ggplot,
    # include the #obs and
    # coefficients in title
    ggplot(.y,aes(x=displ,y=cty))+
      ggtitle(paste0(.x,"(#",nrow(.y),") I=",
                    intercept,
                    " S=",slope))+
      geom_point(color=randomColor(),size=1.5)+
      geom_abline(slope=slope,intercept = intercept)+
      theme_classic()+
      geom_jitter()+
      theme(plot.title = element_text(size = 7,
                                       face = "italic"))
  }))
```

```
mpg1
#> # A tibble: 7 x 5
#> # Groups:   class [7]
#>   class      data          LM      RSquared Plots
#>   <chr>      <list>        <list>    <dbl> <list>
#> 1 suv       <tibble [62 x 10]> <lm>      0.558 <gg>
#> 2 subcompact <tibble [35 x 10]> <lm>      0.527 <gg>
#> 3 pickup     <tibble [33 x 10]> <lm>      0.525 <gg>
#> 4 compact    <tibble [47 x 10]> <lm>      0.358 <gg>
#> 5 midsize    <tibble [41 x 10]> <lm>      0.339 <gg>
#> 6 2seater    <tibble [5 x 10]>  <lm>      0.130 <gg>
#> 7 minivan    <tibble [11 x 10]> <lm>      0.124 <gg>
```


Visualise results (package ggpubr)



```
p1 <- ggarrange(plotlist=pull(mpg1, Plots))  
p1
```